



# PRESENTATION

Presented by: Kislay Kumar



# INTRODUCTION

- Customer churn explained: why customers leave matters
- Importance of predicting churn for business growth
- Data has been taken from kaggle – Telco Customer Churn on Kaggle.

# BACKGROUND

- Problem: Predict if a customer will leave a service (churn) based on usage data
  - Collect & clean customer datasets
- Analyze features like contract type, tenure, payment method with charts
- Build classification models (Decision Trees, Random Forest, or XGBoost)
  - Validate model with accuracy, confusion matrix
    - Generate insights to help reduce churn
    - Document clearly with code and findings

# PROBLEMS

Predict if a customer will leave a service (churn) based on usage data

High churn leads to revenue loss

# GOALS

Predict which customers may churn using data  
- Collect & clean customer datasets

Analyze features like contract type, tenure, payment method with charts

# THEORY

- What is customer churn prediction?
- Common algorithms: Random Forest, Decision Tree, XGBoost
- Key metrics: accuracy, precision, recall, confusion matrix

Customer churn prediction means using data and machine learning to guess which customers might stop using a service soon. It helps businesses keep customers by acting before they leave.



# ANALYSIS

- Dataset overview (Telco Customer Churn)
- Data cleaning steps: handling missing values, encoding categories
  - Visualizations like churn distribution, tenure histogram, correlation heatmap

Visualize data trends like churn rate by contract type, tenure distribution, payment methods.

# RESULT

I Have attached the screenshot of the code and result  
on next page

I Have collected the data from kaggle –  
You can find the dataset here: Telco Customer Churn on  
Kaggle.



## ✓ 1. Problem Definition:

Predict if a customer will churn (leave) based on their data.

```
[ ] import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## ✓ 2. Data Collection & Cleaning:

```
[ ] from google.colab import drive  
    drive.mount('/content/drive')
```

 Mounted at /content/drive

```
[ ] data = pd.read_csv('/content/drive/MyDrive/WA_Fn-UseC_-Telco-Customer-Churn.csv')
```


```
▶ data = data.drop('customerID', axis=1)  
  data['TotalCharges'] = pd.to_numeric(data['TotalCharges'], errors='coerce')  
  data = data.dropna()
```

## ✓ Encode categorical variables


```
[ ] for col in data.select_dtypes(include='object').columns:  
    data[col] = data[col].astype('category').cat.codes
```

### ✓ 3. Data Analysis & Charts:


```
plt.figure(figsize=(6,4))  
sns.countplot(x='Churn', data=data)  
plt.title('Churn Distribution')  
plt.show()
```



```
plt.figure(figsize=(8,5))  
sns.countplot(x='Contract', hue='Churn', data=data)  
plt.title('Churn by Contract Type')  
plt.show()
```




```
plt.figure(figsize=(6,4))  
sns.histplot(data['tenure'], bins=30)  
plt.title('Tenure Distribution')  
plt.show()
```



```
plt.figure(figsize=(10,8))
sns.heatmap(data.corr(), annot=True, fmt=".2f", cmap="coolwarm")
plt.title('Correlation Heatmap')
plt.show()
```



Correlation Heatmap



1.0  
0.8  
0.6  
0.4  
0.2



## ✓ 4. Model Building:

```
[ ] x = data.drop('Churn', axis=1)  
    y = data['Churn']
```

```
[ ] x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
[ ] rf = RandomForestClassifier(random_state=42)  
    rf.fit(x_train, y_train)
```



RandomForestClassifier

RandomForestClassifier(random\_state=42)

## ✓ 5. Model Checking:

```
[ ] y_pred = rf.predict(X_test)
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
    print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

⇒ Accuracy: 0.7848341232227488

Confusion Matrix:

```
[[1384  165]
 [ 289  272]]
```

Classification Report:


	precision	recall	f1-score	support
0	0.83	0.89	0.86	1549
1	0.62	0.48	0.55	561
accuracy			0.78	2110
macro avg	0.72	0.69	0.70	2110
weighted avg	0.77	0.78	0.78	2110

## ✓ 6. Results & Insights:


### Feature importance to identify key churn drivers

```
[ ] importances = rf.feature_importances_  
    feat_names = X.columns  
    feat_imp = pd.Series(importances, index=feat_names).sort_values(ascending=False)
```

```
[ ] plt.figure(figsize=(8,6))  
    sns.barplot(x=feat_imp, y=feat_imp.index)  
    plt.title('Feature Importance')  
    plt.show()
```



```
plt.figure(figsize=(8,6))  
sns.barplot(x=feat_imp, y=feat_imp.index)  
plt.title('Feature Importance')  
plt.show()
```



## ✓ 7. Report & Documentation:

Clearly document each step with comments and save visualizations if needed.

Optional: Hyperparameter tuning for better model

```
[ ] param_grid = {  
    'n_estimators': [100, 200],  
    'max_depth': [None, 10, 20],  
    'min_samples_split': [2, 5]  
}
```

```
[ ] grid_search = GridSearchCV(rf, param_grid, cv=3, scoring='accuracy')  
grid_search.fit(X_train, y_train)  
print("Best params:", grid_search.best_params_)
```

```
➔ Best params: {'max_depth': 10, 'min_samples_split': 2, 'n_estimators': 100}
```

```
[ ] best_rf = grid_search.best_estimator_  
y_pred_best = best_rf.predict(X_test)  
print("Tuned Model Accuracy:", accuracy_score(y_test, y_pred_best))
```

```
➔ Tuned Model Accuracy: 0.7900473933649289
```

THANK YOU

